

# Satellite Image Classification using ConvNet, Keras and TensorFlow

## Introduction

### Problem

Every minute, the world loses an area of forest the size of 48 football fields. And deforestation in the Amazon Basin accounts for the largest share, contributing to reduced biodiversity, habitat loss, climate change, and other devastating effects. But better data about the location of deforestation and human encroachment on forests can help governments and local stakeholders respond more quickly and effectively.

### Class Labels

Datasets comprise of 17 labels.

haze, primary, agriculture, clear, water, habitation, road, cultivation, slash\_burn, cloudy, partly\_cloudy, conventional\_mine, bare\_ground, artisinal\_mine, blooming, selective\_logging, blow\_down



## Neural Networks

In [ ]:

## Convolutional Neural Networks (ConvNet/CNN)

In [ ]:

## Deep Learning

In [ ]:

## TensorFlow

In [ ]:

## Keras

In [ ]:

In [ ]:

## Environment Setup

(Verify these, probably will need to add remove some packages)

## Install

```
conda create -n keras python=3.5 jupyter
activate keras
conda install theano
conda install mingw libpython
pip install tensorflow
pip install keras
pip install scikit-learn
pip install pillow
pip install h5py
pip install tensorflow-gpu
pip install imagenet_utils
activate keras
```

## Open CV

```
conda install -c menpo opencv3
```

## Verify

```
python -c "from keras import backend; print(backend._BACKEND)"
```

## Config

```
python -c "import os; print(os.path.expanduser('~') + '.keras\keras.json')"
```

## Verify GPU

```
python -c "import tensorflow as tf; sess = tf.Session(); hello = tf.constant('Hello, TensorFlow!');
print(sess.run('hello'))"
```

## Run

```
activate keras
```

# Data Preparation

Download and unzip following datasets from Kaggle. Note for this project we'll be using jpg instead of tif. As processing high resolution tif is computationally expensive. With jpg datasets we can achieve satisfactory results, 96% accuracy and top 15% on Kaggle Leader board.

test-jpg.tar.7z

train-jpg.tar.7z

URL: <https://www.kaggle.com/c/planet-understanding-the-amazon-from-space/data>  
(<https://www.kaggle.com/c/planet-understanding-the-amazon-from-space/data>)

```
In [4]: import numpy as np
import pandas as pd
import h5py
import cv2
import scipy.io as sio
import os
import time # Timing
```

**Serialize train and test datasets in hdf5.**

Create onehot encoding for train labels

```

In [5]: test_output_file = 'test-dataset-128.h5'
        train_output_file = "train-dataset-128.h5"

        test_image_path = 'c:/data/amazon/test-jpg'
        train_image_path = 'c:/data/amazon/train-jpg'

        train_max_image_idx = 40478
        test_max_image_idx = 40668

        train_csv = 'c:/data/amazon/train.csv'

        image_resize = (128,128) # Resize images

        x = []

        start_time_data_prep = time.time()

        for i in range(0, test_max_image_idx + 1):
            img = test_image_path + "/test_" + str(i) + ".jpg"
            if i % 5000 == 0:
                print("reading image: {}".format(img))
                img = cv2.imread(img)
                img = cv2.resize(img,image_resize)
                # img = img.transpose((2,0,1))
                x.append(img)

        print('Saving file: {}'.format(test_output_file))
        x = np.array(x)
        f = h5py.File(test_output_file)
        f['x'] = x
        f.close()

        print('Time elapsed: {} seconds'.format(time.time()-start_time_data_prep))

        reading image: c:/data/amazon/test-jpg/test_0.jpg
        reading image: c:/data/amazon/test-jpg/test_5000.jpg
        reading image: c:/data/amazon/test-jpg/test_10000.jpg
        reading image: c:/data/amazon/test-jpg/test_15000.jpg
        reading image: c:/data/amazon/test-jpg/test_20000.jpg
        reading image: c:/data/amazon/test-jpg/test_25000.jpg
        reading image: c:/data/amazon/test-jpg/test_30000.jpg
        reading image: c:/data/amazon/test-jpg/test_35000.jpg
        reading image: c:/data/amazon/test-jpg/test_40000.jpg
        Saving file: test-dataset-128.h5
        Time elapsed: 85.76890587806702 seconds

```

```

In [6]: ## Train dataset
train_output_file = "train-dataset-128.h5"
train_image_path = 'c:/data/amazon/train-jpg'
train_max_image_idx = 40478
train_csv = 'c:/data/amazon/train.csv'

image_resize = (128,128) # Resize images

start_time_data_prep = time.time()

df = pd.read_csv(train_csv)
print('Training dataset shape: {}'.format(df.shape))
df.head()

# Build list with unique labels
label_list = []
for tag_str in df.tags.values:
    labels = tag_str.split(' ')
    for label in labels:
        if label not in label_list:
            label_list.append(label)

print('Labels: {}'.format(label_list))

# Add onehot features for every label
for label in label_list:
    df[label] = df['tags'].apply(lambda x: 1 if label in x.split(' ') else 0)

# Display head
df.head()

y = np.array(df.ix[:,2:])
#print(y.shape)

x = []

for i in range(0, train_max_image_idx + 1):
    img = train_image_path + "/train_" + str(i) + ".jpg"
    if i % 5000 == 0:
        print("reading image: {}".format(img))
    img = cv2.imread(img)
    img = cv2.resize(img,image_resize)
    # img = img.transpose((2,0,1))
    x.append(img)

print('Saving file: {}'.format(train_output_file))
x = np.array(x)
f = h5py.File(train_output_file)
f['x'] = x
f['y'] = y
f.close()

print('Time elapsed: {} seconds'.format(time.time()-start_time_data_prep))

```

```

Training dataset shape: (40479, 2)
Labels: ['haze', 'primary', 'agriculture', 'clear', 'water', 'habitation', 'road', 'cultivation', 'slash_burn', 'cloudy', 'partly_cloudy', 'conventional_mine', 'bare_ground', 'artisinal_mine', 'blooming', 'selective_logging', 'blow_down']
reading image: c:/data/amazon/train-jpg/train_0.jpg
reading image: c:/data/amazon/train-jpg/train_5000.jpg
reading image: c:/data/amazon/train-jpg/train_10000.jpg
reading image: c:/data/amazon/train-jpg/train_15000.jpg
reading image: c:/data/amazon/train-jpg/train_20000.jpg
reading image: c:/data/amazon/train-jpg/train_25000.jpg
reading image: c:/data/amazon/train-jpg/train_30000.jpg
reading image: c:/data/amazon/train-jpg/train_35000.jpg
reading image: c:/data/amazon/train-jpg/train_40000.jpg
Saving file: train-dataset-128.h5
Time elapsed: 102.22684693336487 seconds

```

## One hot encoding for the labels

In [8]: df.head()

Out[8]:

	image_name	tags	haze	primary	agriculture	clear	water	habitation	road	cul
0	train_0	haze primary	1	1	0	0	0	0	0	0
1	train_1	agriculture clear primary water	0	1	1	1	1	0	0	0
2	train_2	clear primary	0	1	0	1	0	0	0	0
3	train_3	clear primary	0	1	0	1	0	0	0	0
4	train_4	agriculture clear habitation primary road	0	1	1	1	0	1	1	0

## Model

## Description

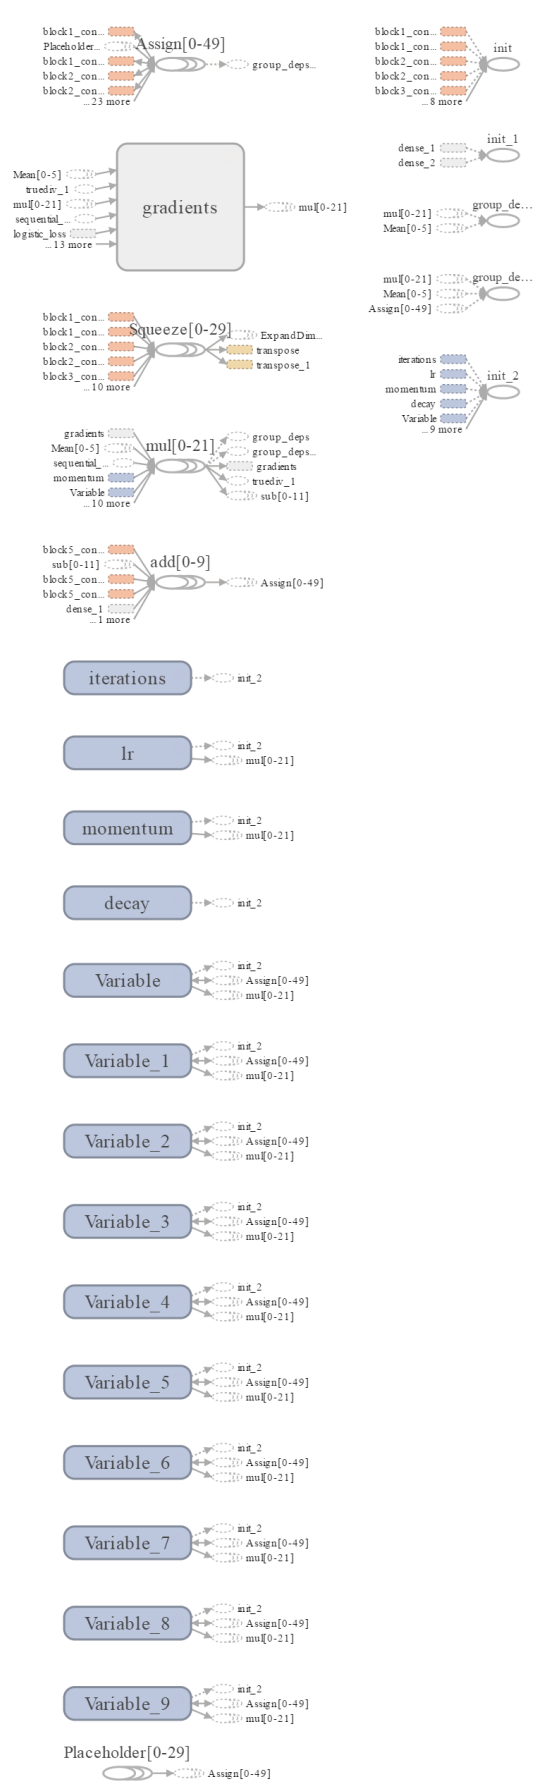
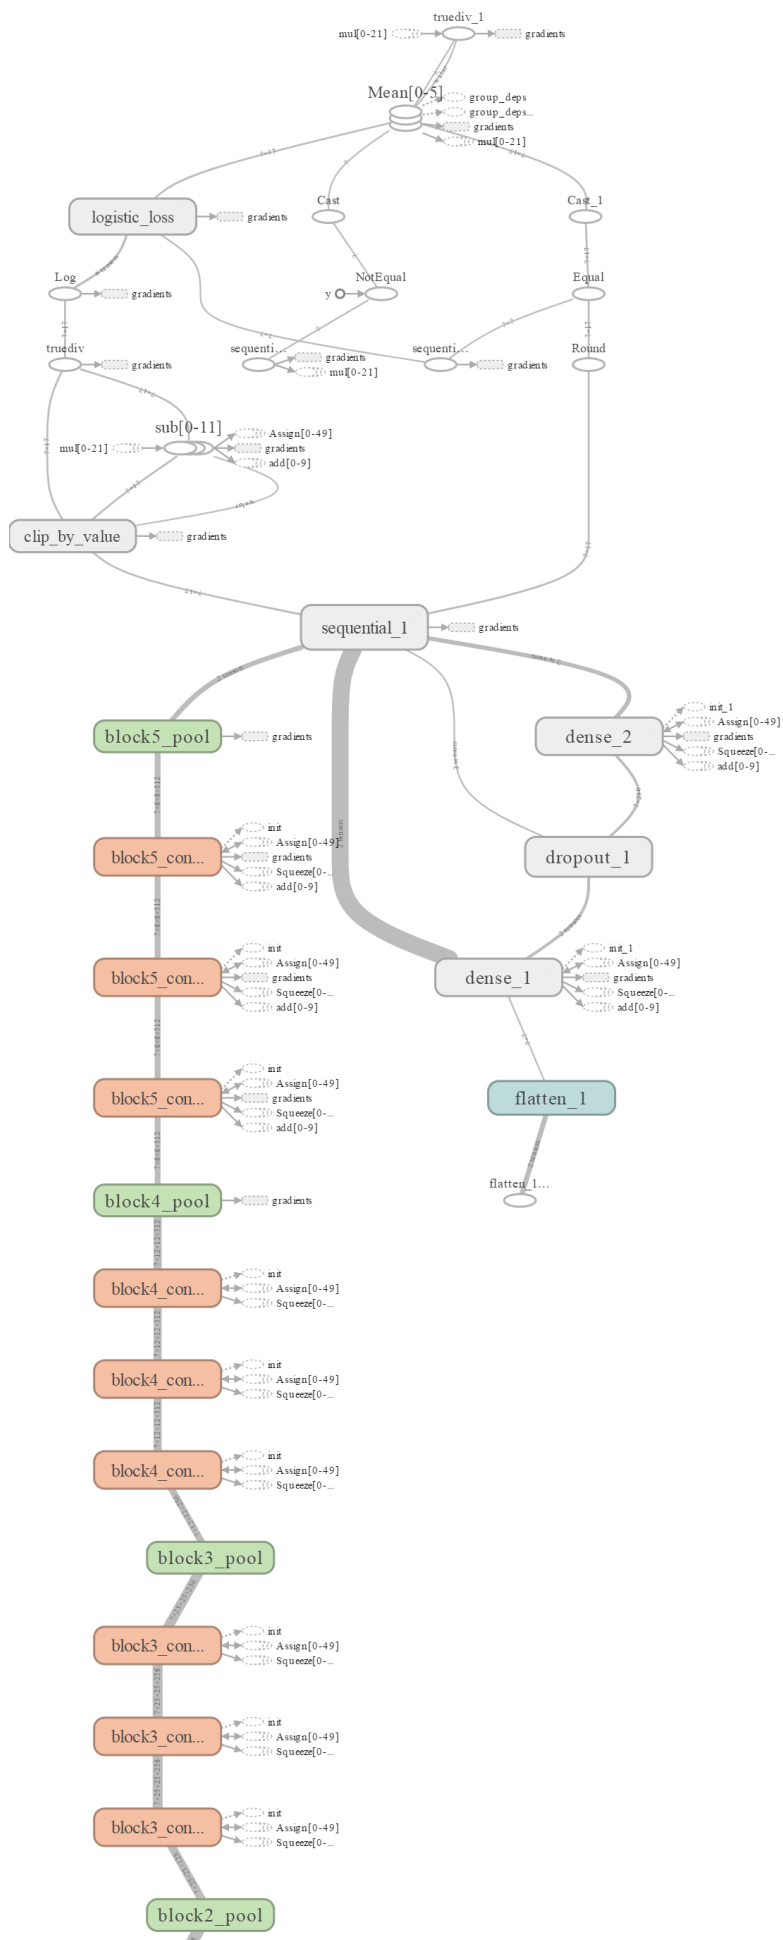
In [ ]:

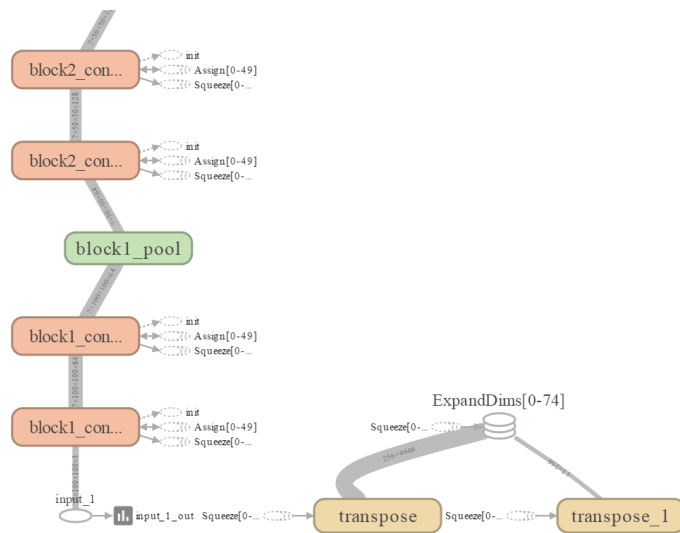
### TensorBoard Graph

VGG16 pre-trained model frozen upto conv4. Set the first 15 layers (up to the conv4) to non-trainable (weights will not be updated)









In [ ]:

```
In [1]: import numpy as np
from keras.models import Sequential, Model
from keras.layers import Dense, Dropout, Activation, Flatten, Reshape
from keras.layers import Conv2D, MaxPooling2D
from keras.callbacks import ModelCheckpoint, TensorBoard
from keras import applications
from keras.optimizers import adam
import h5py
from sklearn.model_selection import train_test_split
import time
```

Using TensorFlow backend.

## Function to load training dataset for a given batch size and train/validation split

```
In [2]: def load_train_dataset(dataset, random_state, batch_size=-1, test_size=0.2):
    """
    batch_size=-1 returns all
    """
    f = h5py.File(dataset)
    if batch_size == -1:
        x = f['x'].value
        y = f['y'].value
    else:
        x = f['x'][:batch_size,]
        y = f['y'][:batch_size,]

    f.close()
    x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=test_size,
        random_state=random_state)

    return x_train, x_test, y_train, y_test
```

## Function to load test dataset

```
In [3]: def load_test_dataset(dataset):  
        f = h5py.File(dataset)  
        x = f['x'].value    #[0:100,]  
        f.close()  
  
        return x
```

```
In [ ]:
```

## Save Bottleneck Features from VGG16 Model

### Global settings

```
In [4]: epochs = 3  
        batch_size = 8  
        bottleneck_features_train_path = 'bottleneck_features_train_128_vgg.npy'  
        bottleneck_features_validation_path = 'bottleneck_features_validation_128_vgg.npy'  
        # path to the model weights files.  
        top_model_weights_path = 'bottleneck_fc_model_128_vgg.h5'  
        # dimensions of our images.  
        input_shape = (128,128,3)
```

### Save Bottleneck Features

```
In [5]: def save_bottleneck_features():  
        # build the VGG16 network  
        # First time it will take longer, as it downloads the weights.  
        model = applications.VGG16(include_top=False, weights='imagenet')  
        model.summary()  
        start_time = time.time()  
  
        bottleneck_features_train = model.predict(  
            x_train, batch_size = batch_size, verbose=1)  
  
        np.save(open(bottleneck_features_train_path, 'wb'),  
            bottleneck_features_train)  
  
        bottleneck_features_validation = model.predict(  
            x_test, batch_size = batch_size, verbose=1)  
        np.save(open(bottleneck_features_validation_path, 'wb'),  
            bottleneck_features_validation)  
        print('save_bottleneck_features(): Time elapsed: {} seconds'.format(time.time()-start_time))
```

## Train Top Model

Train fully connected model.

```
In [19]: def train_top_model():
    start_time = time.time()
    train_data = np.load(open(bottleneck_features_train_path, 'rb'))
    train_labels = y_train
    validation_data = np.load(open(bottleneck_features_validation_path, 'rb'))
    validation_labels = y_test

    print('train_data.shape[1:]: {}'.format(train_data.shape[1:]))
    model = Sequential()
    model.add(Flatten(input_shape=train_data.shape[1:]))
    model.add(Dense(256, activation='relu'))
    model.add(Dropout(0.5))
    model.add(Dense(17, activation='sigmoid'))

    model.compile(optimizer='rmsprop',
                  loss='binary_crossentropy', metrics=['accuracy'])

    ##### Update this if using weights from previous run #####
    # model.load_weights(top_model_weights_path)
    model.fit(train_data, train_labels,
              epochs=epochs,
              batch_size=batch_size,
              validation_data=(validation_data, validation_labels))
    model.save_weights(top_model_weights_path)
    print('train_top_model(): Time elapsed: {} seconds'.format(time.time()-start_time_data_prep))
```

Putting them together and train full model

```
In [12]: def train_full_model():
    start_time = time.time()
    # build the VGG16 network
    base_model = applications.VGG16(weights='imagenet', include_top=False, input_shape=input_shape)
    print('Model loaded.')
    # base_model.summary()
    # build a classifier model to put on top of the convolutional model
    top_model = Sequential()
    top_model.add(Flatten(input_shape=(4,4,512)))
    top_model.add(Dense(256, activation='relu'))
    top_model.add(Dropout(0.5))
    top_model.add(Dense(17, activation='sigmoid'))

    # note that it is necessary to start with a fully-trained
    # classifier, including the top classifier,
    # in order to successfully do fine-tuning
    top_model.load_weights(top_model_weights_path)
```

```

# top_model.summary()

# add the model on top of the convolutional base
model = Model(inputs= base_model.input, outputs= top_model(base_model.output))

# set the first 15 layers (up to the last conv block)
# to non-trainable (weights will not be updated)
for layer in model.layers[:15]:
    # print(layer.name)
    layer.trainable = False

model.summary()

#### Update following if using weights from previous run ####
#model.load_weights('weights-model-07.01-0.95972.hdf5')

# compile the model with adam optimizer
# and a very slow learning rate.
model.compile(loss='binary_crossentropy',
              optimizer=adam(lr=1e-4),
              metrics=['accuracy'])

x_train, x_test, y_train, y_test = load_train_dataset(random_state=random_
state, dataset=train_dataset, test_size=test_size)

x_train = x_train.astype('float32')
x_test  = x_test.astype('float32')

x_train /= 255
x_test  /= 255

tbCallBack = TensorBoard(log_dir='graph', histogram_freq=0, write_graph=True,
write_images=False, embeddings_freq=0)
check = ModelCheckpoint("weights-model-07.{epoch:02d}-{val_acc:.5f}.hdf5",
monitor='val_acc', verbose=1,
                        save_best_only=True, save_weights_only=True, mode='auto')
model.fit(x_train, y_train, batch_size=batch_size,
epochs=epochs, callbacks=[check], validation_data=(x_test, y_test))

## Predict
print('Generate Predictions...')
x_test = load_test_dataset(test_dataset)
x_test = x_test.astype('float32')
x_test /= 255.

best_threshold = [0.2] * 17

# print("best_threshold: {}".format(best_threshold))

pred = model.predict(x_test, verbose=1, batch_size=8)
# print(pred)
print(pred.shape)

classes = ['haze',

```

```

        'primary',
        'agriculture',
        'clear',
        'water',
        'habitation',
        'road',
        'cultivation',
        'slash_burn',
        'cloudy',
        'partly_cloudy',
        'conventional_mine',
        'bare_ground',
        'artisinal_mine',
        'blooming',
        'selective_logging',
        'blow_down']

y_pred = []

##text=List of strings to be written to file
with open(submission_file, 'w') as file:
    file.write("image_name,tags")
    file.write('\n')

    for i in range(pred.shape[0]):
        y_pred = np.array([1 if pred[i, j] >= best_threshold[j] else 0 for
j in range(pred.shape[1])])
        # print(y_pred)

        # extracting actual class name
        y_pred = [classes[i] for i in range(17) if y_pred[i] == 1]
        y_pred = " ".join([str(item) for item in y_pred])
        # print(y_pred)
        line = "test_{},{}".format(i, y_pred)
        file.write(line)
        file.write('\n')

    print('train_full_model(): Time elapsed: {} seconds'.format(time.time()-start_time))

```

## Execution



```
In [8]: random_state = 101
train_dataset = 'train-dataset-128.h5'
test_size = 0.3 # for the train/test split
submission_file = 'model-07-submission.csv'

x_train, x_test, y_train, y_test = load_train_dataset(dataset=train_dataset,
                                                        random_state=random_state,
                                                        test_size=test_size)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

x_train /= 255
x_test /= 255
```

**Save bottleneck features to npy arrays**

```
In [9]: save_bottleneck_features()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, None, None, 3)	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0

=====  
Total params: 14,714,688  
Trainable params: 14,714,688  
Non-trainable params: 0  
=====  
28335/28335 [=====] - 1269s  
12144/12144 [=====] - 540s  
save\_bottleneck\_features(): Time elapsed: 1816.2711477279663 seconds

## Train Top Model with the Bottleneck Features

```
In [18]: # Ideally you should run 10 epochs
# If you are re-running, make sure to load the
# weights from last run
epochs = 1
# Adjust this per your HW (from 8-32)
batch_size = 8

train_top_model()

train_data.shape[1:]: (4, 4, 512)
Train on 28335 samples, validate on 12144 samples
Epoch 1/1
28335/28335 [=====] - 65s - loss: 0.2107 - acc: 0.92
12 - val_loss: 0.1816 - val_acc: 0.9319
train_top_model(): Time elapsed: 77.15817737579346 seconds
```

## Train Full Model

```
In [21]: # Ideally you should run 10 epochs
# If you are re-running, make sure to load the
# weights from last run
epochs = 1
# Adjust this per your HW (from 8-32)
batch_size = 8

random_state = 55
train_dataset = 'train-dataset-128.h5'
test_dataset = 'test-dataset-128.h5'
test_size = 0.3 # for the train/test split
submission_file = 'model-07-submission.csv'

train_full_model()
```

Model loaded.

Layer (type)	Output Shape	Param #
input_5 (InputLayer)	(None, 128, 128, 3)	0
block1_conv1 (Conv2D)	(None, 128, 128, 64)	1792
block1_conv2 (Conv2D)	(None, 128, 128, 64)	36928
block1_pool (MaxPooling2D)	(None, 64, 64, 64)	0
block2_conv1 (Conv2D)	(None, 64, 64, 128)	73856
block2_conv2 (Conv2D)	(None, 64, 64, 128)	147584
block2_pool (MaxPooling2D)	(None, 32, 32, 128)	0
block3_conv1 (Conv2D)	(None, 32, 32, 256)	295168
block3_conv2 (Conv2D)	(None, 32, 32, 256)	590080
block3_conv3 (Conv2D)	(None, 32, 32, 256)	590080
block3_pool (MaxPooling2D)	(None, 16, 16, 256)	0
block4_conv1 (Conv2D)	(None, 16, 16, 512)	1180160
block4_conv2 (Conv2D)	(None, 16, 16, 512)	2359808
block4_conv3 (Conv2D)	(None, 16, 16, 512)	2359808
block4_pool (MaxPooling2D)	(None, 8, 8, 512)	0
block5_conv1 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv2 (Conv2D)	(None, 8, 8, 512)	2359808
block5_conv3 (Conv2D)	(None, 8, 8, 512)	2359808
block5_pool (MaxPooling2D)	(None, 4, 4, 512)	0
sequential_5 (Sequential)	(None, 17)	2101777
Total params: 16,816,465		
Trainable params: 9,181,201		
Non-trainable params: 7,635,264		

Train on 28335 samples, validate on 12144 samples

Epoch 1/1

28328/28335 [=====>.] - ETA: 0s - loss: 0.2339 - acc: 0.9060  
Epoch 00000: val\_acc improved from -inf to 0.90660, saving model to weights-model-07.00-0.90660.hdf5

28335/28335 [=====] - 2296s - loss: 0.2339 - acc: 0.9060 - val\_loss: 0.2071 - val\_acc: 0.9066

Predictions...

40669/40669 [=====] - 1836s